

Average Processing Time per Machine

Objective:

Calculate the average time each machine takes to complete a process based on start and end timestamps.

Tables Involved:

- **Activity**
 - machine_id (int): Identifier for the machine.
 - process_id (int): Identifier for the process.
 - activity_type (enum): Type of activity ('start' or 'end').
 - timestamp (float): Time at which the activity occurred.

Approach:

1. Pairing Start and End Timestamps:

- Use a self-join on the Activity table to pair each 'start' activity with its corresponding 'end' activity based on machine_id and process_id.

2. Calculating Processing Time:

- For each pair, compute the difference between the 'end' and 'start' timestamps to determine the processing time for each process.

3. Averaging Processing Times:

- Group the results by machine_id and calculate the average processing time for each machine.

4. Rounding:

- Round the average processing time to three decimal places for precision.

SQL Query:

SELECT

start.machine_id,

ROUND(AVG(end.timestamp - start.timestamp), 3) AS processing_time

FROM

Activity AS start

JOIN Activity AS end

ON start.machine_id = end.machine_id

AND start.process_id = end.process_id

AND start.activity_type = 'start'

AND end.activity_type = 'end'

GROUP BY

start.machine_id;

• Case 1

Input

Activity =

machine_id	process_id	activity_type	timestamp
0	0	start	0.712
0	0	end	1.52
0	1	start	3.14
0	1	end	4.12
1	0	start	0.55
1	0	end	1.55
1	1	start	0.43
1	1	end	1.42
2	0	start	4.1
2	0	end	4.512
2	1	start	2.5
2	1	end	5

[View less](#)

Output

machine_id	processing_time
0	0.894
1	0.995
2	1.456

Expected

machine_id	processing_time
0	0.894
1	0.995
2	1.456

Employee Bonus

Objective:

Retrieve the names and bonus amounts of employees who have a bonus less than 1000 or no bonus recorded.

Tables Involved:

1. Employee

- empld (Primary Key): Unique identifier for each employee.
- name: Name of the employee.
- supervisor: ID of the employee's supervisor.
- salary: Salary of the employee.

2. Bonus

- empld (Primary Key, Foreign Key referencing Employee.empld): Unique identifier linking to the Employee table.
- bonus: Bonus amount received by the employee.

Approach:

- Perform a LEFT JOIN between the Employee and Bonus tables on the empld column to include all employees, regardless of whether they have a corresponding bonus record. [CodingTutorials | Codelabs365](#)
- Filter the results to include only those employees whose bonus is less than 1000 or who do not have a bonus record (i.e., bonus is NULL).

SQL Query:

```
SELECT e.name, b.bonus
FROM Employee e
LEFT JOIN Bonus b ON e.empld = b.empld
WHERE b.bonus < 1000 OR b.bonus IS NULL;
```

- Case 1

Input

Employee =

empId	name	supervisor	salary
3	Brad	null	4000
1	John	3	1000
2	Dan	3	2000
4	Thomas	3	4000

Bonus =

empId	bonus
2	500
4	2000

Output

name	bonus
Brad	null
John	null
Dan	500

Expected

name	bonus
Brad	null
John	null
Dan	500

Counting Student Exam Attendances

Objective:

Determine the number of times each student attended each exam.

Tables Involved:

1. Students

- student_id (Primary key)
- student_name

2. Subjects

- subject_name (Primary key)

3. Examinations

- student_id
- subject_name

Approach:

1. Generate All Student-Subject Combinations:

- Perform a CROSS JOIN between the Students and Subjects tables to create all possible student-subject pairs. This ensures that even if a student hasn't attended an exam for a subject, the combination is still represented.

2. Count Exam Attendances:

- Use a LEFT JOIN with the Examinations table on the student_id and subject_name to count the occurrences of each student attending each exam. The COUNT function will tally the number of times each student has attended the exam for a particular subject.

3. Group and Order Results:

- Group the results by student_id, student_name, and subject_name to aggregate the attendance counts.
- Order the final output by student_id and subject_name for clarity.

SQL Query:

SELECT

s.student_id,

s.student_name,

```
    sub.subject_name,  
    COUNT(e.student_id) AS attended_exams  
FROM  
    Students s  
CROSS JOIN  
    Subjects sub  
LEFT JOIN  
    Examinations e  
ON  
    s.student_id = e.student_id  
    AND sub.subject_name = e.subject_name  
GROUP BY  
    s.student_id,  
    s.student_name,  
    sub.subject_name  
ORDER BY  
    s.student_id,  
    sub.subject_name;
```

Input

Students =

student_id	student_name
1	Alice
2	Bob
13	John
6	Alex

Subjects =

subject_name
Math
Physics
Programming

Examinations =

student_id	subject_name
1	Math
1	Physics
1	Programming
2	Programming
1	Physics
1	Math
13	Math
13	Programming
13	Physics
2	Math
1	Math

[View less](#)

Output

student_id	student_name	subject_name	attended_exams
1	Alice	Math	3
1	Alice	Physics	2
1	Alice	Programming	1
2	Bob	Math	1
2	Bob	Physics	0
2	Bob	Programming	1
6	Alex	Math	0
6	Alex	Physics	0
6	Alex	Programming	0
13	John	Math	1
13	John	Physics	1
13	John	Programming	1

[View less](#)

Expected

student_id	student_name	subject_name	attended_exams
1	Alice	Math	3
1	Alice	Physics	2
1	Alice	Programming	1
2	Bob	Math	1
2	Bob	Physics	0
2	Bob	Programming	1
6	Alex	Math	0
6	Alex	Physics	0
6	Alex	Programming	0
13	John	Math	1
13	John	Physics	1
13	John	Programming	1

[View less](#)

Managers with at Least 5 Direct Reports

Objective:

Identify managers who have five or more direct reports.

Table Involved:

- **Employee**
 - id (Primary key)
 - name (Employee's name)
 - department (Department name)
 - managerId (ID of the employee's manager)

Approach:

1. Self-Join Method:

- Perform a self-join on the Employee table:[LinkedIn](#)
 - Join Employee as e with Employee as m on e.managerId = m.id.
- Group the results by m.id and m.name.
- Use the HAVING clause to filter groups where the count of e.id is 5 or more.

SQL Query:

```
SELECT m.name
FROM Employee e
JOIN Employee m ON e.managerId = m.id
GROUP BY m.id, m.name
HAVING COUNT(e.id) >= 5;
```

2. Subquery with WHERE Clause:

- Use a subquery to find managerIds with at least five direct reports:[DEV Community](#)
 - Select managerId from Employee, group by managerId, and filter with HAVING COUNT(*) >= 5.
- In the main query, select name from Employee where id is in the list of managerIds obtained from the subquery.

3.

SQL Query:

```
SELECT name
FROM Employee
WHERE id IN (
    SELECT managerId
    FROM Employee
    GROUP BY managerId
    HAVING COUNT(*) >= 5
);
```

3. Common Table Expression (CTE):

- Define a CTE to find managerIds with at least five direct reports:
 - Select managerId from Employee, group by managerId, and filter with HAVING COUNT(*) >= 5.
- In the main query, select name from Employee where id matches managerId from the CTE. [Medium](#)

SQL Query:

```
WITH ManagerCounts AS (
    SELECT managerId
    FROM Employee
    GROUP BY managerId
    HAVING COUNT(*) >= 5
)
SELECT name
FROM Employee
WHERE id IN (SELECT managerId FROM ManagerCounts);
```

• Case 1

Input

Employee =

id	name	department	managerId
---	----	-----	-----
101	John	A	null
102	Dan	A	101
103	James	A	101
104	Amy	A	101
105	Anne	A	101
106	Ron	B	101

Output

name

John

Expected

name

John